

6

Linear Algebra Least Squares

Anahita Zarei, Ph.D.

Overview

- Dot Product
- Distance
- Unit Vector
- Orthogonality
- Least squares solution to inconsistent systems
- Reading: Chapter 6 with emphasis on sections 6.1, 6.5, 6.6

Section 6.1

INNER PRODUCT, LENGTH, AND ORTHOGONALITY

INNER PRODUCT

- If \mathbf{u} and \mathbf{v} are vectors in \mathbb{R}^n , then we regard \mathbf{u} and \mathbf{v} as $n \times 1$ matrices.

$$\mathbf{u} = \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix} \quad \mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}$$

- The inner product (dot product) of \mathbf{u} and \mathbf{v} is
 $u_1 v_1 + u_2 v_2 + \cdots + u_n v_n.$

INNER PRODUCT

- Can you express the dot product in terms of matrix multiplication of \mathbf{u} and \mathbf{v} ?
- Hint: The transpose \mathbf{u}^T is a $1 \times n$ matrix, and the matrix product $\mathbf{u}^T \mathbf{v}$ is a 1×1 matrix, which we write as a single real number (a scalar) without brackets.

INNER PRODUCT

■ If $\mathbf{u} = \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}$,

then the inner product of \mathbf{u} and \mathbf{v} is

$$\begin{bmatrix} u_1 & u_2 & \cdots & u_n \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = u_1 v_1 + u_2 v_2 + \cdots + u_n v_n \quad .$$

Dot Product in R - Example

- $u = \begin{bmatrix} 1 \\ -3 \\ 4 \end{bmatrix} \quad v = \begin{bmatrix} 2 \\ 1 \\ -2 \end{bmatrix}$
- Find the dot product of u and v in R.

```
> u = matrix(c(1,-3,4), nrow = 3)
> v = matrix(c(2,1,-2), nrow = 3)
```

```
> t(u)%*%v
      [,1]
[1,]    -9
```

```
> library(pracma)
> dot(u,v)
[1] -9
```

Dot Product in Python - Example

- $u = \begin{bmatrix} 1 \\ -3 \\ 4 \end{bmatrix}$ $v = \begin{bmatrix} 2 \\ 1 \\ -2 \end{bmatrix}$
- Find the dot product of u and v in Python.

```
import numpy as np
```

```
u = np.array([1, -3, 4])
```

```
v = np.array([2, 1, -2])
```

```
np.dot(u, v)
```

-9

Arrays in Python

- Note that when you enter array U in Python as follows it'll be a 1D array by default.

```
u = np.array([1,-3,4])  
print(u)  
print(u.shape)
```

```
[ 1 -3  4]  
(3,)
```

- You can make it a row or a column vector by using the reshape command.

```
: rowU = u.reshape(1,3)  
print(rowU)  
print(rowU.shape)
```

```
[[ 1 -3  4]]  
(1, 3)
```

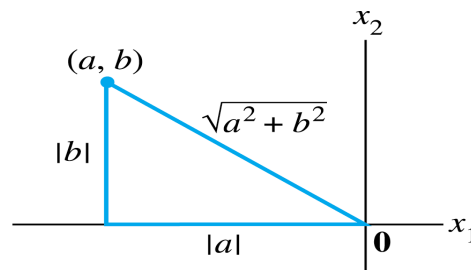
- You can also perform the dot product using matrix multiplication.

```
In [28]: colV = v.reshape(3,1)  
print(rowU@colV)
```

```
[[ -9]]
```

THE LENGTH OF A VECTOR

- If we identify $\mathbf{v} = \begin{bmatrix} a \\ b \end{bmatrix}$ with a geometric point in the plane, then $\|\mathbf{v}\|$ coincides with the standard notion of the length of the line segment from the origin to \mathbf{v} .
- This follows from the Pythagorean Theorem applied to a triangle such as the one shown in the following figure.



Interpretation of $\|\mathbf{v}\|$ as length.

THE LENGTH OF A VECTOR

- **Definition:** The **length** (or **norm**) of \mathbf{v} is the nonnegative scalar $\|\mathbf{v}\|$ defined by

$$\|\mathbf{v}\| = \sqrt{v_1^2 + v_2^2 + \cdots + v_n^2}$$

- Can you express norm of \mathbf{v} in terms of a dot product?
- $\|\mathbf{v}\| = \sqrt{\mathbf{v} \cdot \mathbf{v}}$

THE LENGTH OF A VECTOR in R

- Find length of $u = \begin{bmatrix} 1 \\ -3 \\ 4 \end{bmatrix}$ in R.

```
> sqrt(t(u)%*%u)
      [,1]
[1,] 5.09902
```

```
> norm(u, '2')
[1] 5.09902
```

- Note that R by default finds L1 norm, unless you specify what type of norm needs to be calculated.

THE LENGTH OF A VECTOR in Python

- Find length of $u = \begin{bmatrix} 1 \\ -3 \\ 4 \end{bmatrix}$ in Python.

```
import math  
math.sqrt(np.dot(u,u))
```

5.0990195135927845

```
from numpy.linalg import norm  
norm(u)
```

5.0990195135927845

Alternative method for finding the dot product

- If \mathbf{u} and \mathbf{v} are nonzero vectors in either \mathbb{R}^2 or \mathbb{R}^3 , then there is a nice connection between their inner product and the angle between the two line segments from the origin to the points identified with \mathbf{u} and \mathbf{v} .
- The formula is $\mathbf{u} \cdot \mathbf{v} = \|\mathbf{u}\| \|\mathbf{v}\| \cos \theta$ (proof on the next slide)
- Two vectors \mathbf{u} and \mathbf{v} are **orthogonal** (to each other) if $\mathbf{u} \cdot \mathbf{v} = 0$.

Proof

Unit Vector

- A vector whose length is 1 is called a **unit vector**.
- If we *divide* a nonzero vector \mathbf{v} by its length—that is, multiply by $1 / \|\mathbf{v}\|$ —we obtain a unit vector \mathbf{u} because the length of \mathbf{u} is $(1 / \|\mathbf{v}\|) \|\mathbf{v}\|$.
- The process of creating \mathbf{u} from \mathbf{v} is sometimes called **normalizing** \mathbf{v} , and we say that \mathbf{u} is *in the same direction* as \mathbf{v} .

Finding Unit Vector – Example R

- **Example 2:** Let $\mathbf{v} = (1, -2, 2, 0)$. (Assume \mathbf{v} is a column vector). Find a unit vector \mathbf{u} in the same direction as \mathbf{v} .
`> v = matrix(c(1,-2,2,0), nrow = 4)`

- **Solution:** First, compute the length of \mathbf{v} :

$$\|\mathbf{v}\|^2 = \mathbf{v}^T \mathbf{v} = (1)^2 + (-2)^2 + (2)^2 + (0)^2 = 9$$

$$\|\mathbf{v}\| = \sqrt{9} = 3$$

```
> norm(v, '2')  
[1] 3
```

- Then, multiply \mathbf{v} by $1 / \|\mathbf{v}\|$ to obtain

$$\mathbf{u} = \frac{1}{\|\mathbf{v}\|} \mathbf{v} = \frac{1}{3} \mathbf{v} = \frac{1}{3} \begin{bmatrix} 1 \\ -2 \\ 2 \\ 0 \end{bmatrix} = \begin{bmatrix} 1/3 \\ -2/3 \\ 2/3 \\ 0 \end{bmatrix}$$

```
> u=v/norm(v, '2')  
> print(u)  
           [,1]  
[1,] 0.3333333  
[2,] -0.6666667  
[3,] 0.6666667  
[4,] 0.0000000
```

Finding Unit Vector – Example Python

- **Example 2:** Let $\mathbf{v} = (1, -2, 2, 0)$. (Assume \mathbf{v} is a column vector). Find a unit vector \mathbf{u} in the same direction as \mathbf{v} .

```
a = np.array([1, -2, 2, 0])  
unitA = a/norm(a)  
print(unitA)
```

```
[ 0.33333333 -0.66666667  0.66666667  0.]
```

Verify your answer

- To check that $\|u\| = 1$, it suffices to show that $\|u\|^2 = 1$

$$\begin{aligned}\|u\|^2 &= u \cdot u = \left(\frac{1}{3}\right)^2 + \left(-\frac{2}{3}\right)^2 + \left(\frac{2}{3}\right)^2 + (0)^2 \\ &= \frac{1}{9} + \frac{4}{9} + \frac{4}{9} + 0 = 1\end{aligned}$$

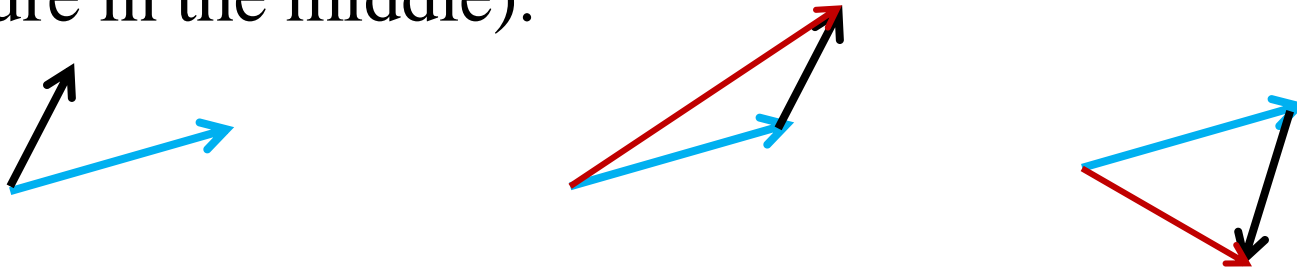
```
> norm(u, '2')  
[1] 1
```

```
In [44]: norm(unitA)
```

```
Out[44]: 1.0
```

DISTANCE IN \mathbb{R}^n

- **Definition:** For \mathbf{u} and \mathbf{v} in \mathbb{R}^n , the **distance between \mathbf{u} and \mathbf{v}** , written as $\text{dist}(\mathbf{u}, \mathbf{v})$, is the length of the vector $\mathbf{u} - \mathbf{v}$. That is, $\text{dist}(\mathbf{u}, \mathbf{v}) = \|\mathbf{u} - \mathbf{v}\|$
- Recall that we can geometrically add two vectors by moving the tail of one vector to the head of the other (figure in the middle).



- Subtraction can be done similarly by adding the negative of the vector. (figure on the right)

Example

- **Example 4:** Compute the distance between the vectors $u = (7,1)$ and $v = (3,2)$

Solution

- **Solution:** Calculate

```
: u = np.array([7,1])  
v = np.array([3,2])  
norm(u-v)  
: 4.123105625617661
```

$$\mathbf{u} - \mathbf{v} = \begin{bmatrix} 7 \\ 1 \end{bmatrix} - \begin{bmatrix} 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 4 \\ -1 \end{bmatrix}$$

$$\|\mathbf{u} - \mathbf{v}\| = \sqrt{4^2 + (-1)^2} = \sqrt{17}$$

Section 6.5

Least Squares Problems

Background

- Consider $Ax=b$ where the system is inconsistent. i.e. there doesn't exist a solution x that satisfies this matrix equation.
- When a solution is demanded and none exists, the best one can do is to find an x that makes **Ax as close as possible to b .**
- Think of Ax as an approximation to b . The smaller the distance between b and Ax given by $\|b-Ax\|$, the better the approximation.

LEAST-SQUARES PROBLEMS

- **Definition:** If A is $m \times n$ and \mathbf{b} is in \mathbb{R}^m , a **least-squares solution** of $A\mathbf{x} = \mathbf{b}$ is an $\hat{\mathbf{x}}$ in \mathbb{R}^n such that

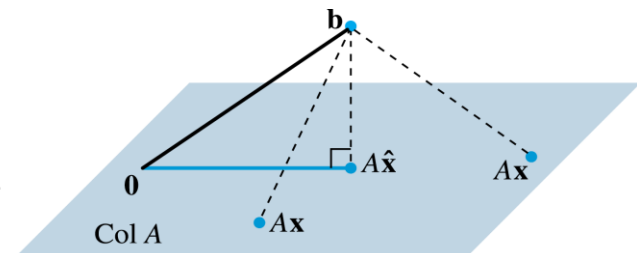
$$\|\mathbf{b} - A\hat{\mathbf{x}}\| \leq \|\mathbf{b} - A\mathbf{x}\|$$

for all \mathbf{x} in \mathbb{R}^n .

- Note that vector $A\mathbf{x}$ is a linear combinations of columns of A . (Recall the 2nd version of matrix-vector multiplication definition).
- This implies that no matter what \mathbf{x} we select, the vector $A\mathbf{x}$ will be on the plane spanned by columns of A (column space of A)

Shortest Distance

- So we seek an x that makes Ax the closest point in Col A to b .
- Recall that the shortest distance is the orthogonal distance.
- Therefore, the dotted line that is perpendicular to the plane will result in smallest distance between b and Ax .
- What's the expression for the dotted line?



The vector b is closer to $A\hat{x}$ than to Ax for other x .

Shortest Distance

- If a vector is orthogonal to a plane it'll be orthogonal to any vector on that plane, including columns of matrix A , namely a_1 and a_2 .
- So $b - A\hat{x}$ is orthogonal to each column of A .
- If \mathbf{a}_j is any column of A , then $a_j \cdot (b - A\hat{x}) = 0$, and $a_j^T (b - A\hat{x}) = 0$.

SOLUTION OF THE GENREAL LEAST-SQUARES PROBLEM

- Since each \mathbf{a}_j^T is a row of A^T ,
$$A^T (\mathbf{b} - A\hat{\mathbf{x}}) = 0 \quad (2)$$

- Thus
$$A^T \mathbf{b} - A^T A \hat{\mathbf{x}} = 0$$

$$A^T A \hat{\mathbf{x}} = A^T \mathbf{b}$$

- These calculations show that each least-squares solution of $A\mathbf{x} = \mathbf{b}$ satisfies the equation

$$A^T A \mathbf{x} = A^T \mathbf{b} \quad (3)$$

- The matrix equation (3) represents a system of equations called the **normal equations** for $A\mathbf{x} = \mathbf{b}$.
- A solution of (3) is often denoted by $\hat{\mathbf{x}}$.

$$\hat{\mathbf{x}} = (A^T A)^{-1} A^T \mathbf{b}$$

SOLUTION OF THE GENREAL LEAST-SQUARES PROBLEM

- **Example 1:** Find a least-squares solution of the inconsistent system $Ax = b$ for

$$A = \begin{bmatrix} 4 & 0 \\ 0 & 2 \\ 1 & 1 \end{bmatrix}, b = \begin{bmatrix} 2 \\ 0 \\ 11 \end{bmatrix}$$

Solution in R

```
> A = matrix(c(4,0,0,2,1,1), nrow = 3, byrow = TRUE)
> b = matrix(c(2,0,11), nrow = 3)
> x = inv(t(A)%*%A)%*%t(A)%*%b
> print(x)
```

	[,1]
[1,]	1
[2,]	2

OR

```
> library(MASS)
> x = ginv(A)%*%b
> print(x)
```

	[,1]
[1,]	1
[2,]	2

Solution in Python

```
: A = np.array([[4,0],[0,2],[1,1]])  
b = np.array([[2],[0],[11]])  
print(A)  
print(b)
```

```
[[4 0]  
 [0 2]  
 [1 1]]  
[[ 2]  
 [ 0]  
 [11]]
```

OR

```
: from numpy.linalg import pinv  
x = pinv(A)@b  
print(x)
```

```
[[1.]  
 [2.]]
```

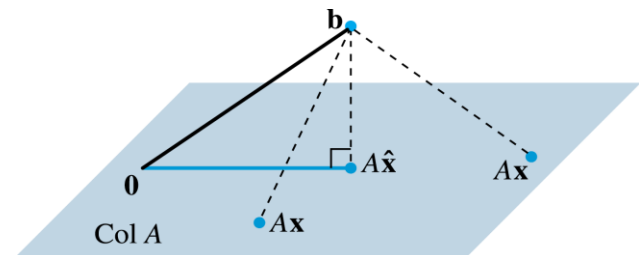
```
from numpy.linalg import inv  
x = inv(A.T@A)@A.T@b  
print(x)
```

```
[[1.]  
 [2.]]
```

Least Squares Error in R and Python

- Determine the least-squares error in previous example.

```
> err = norm(A%*%x - b, '2')  
> print(err)  
[1] 9.165151
```

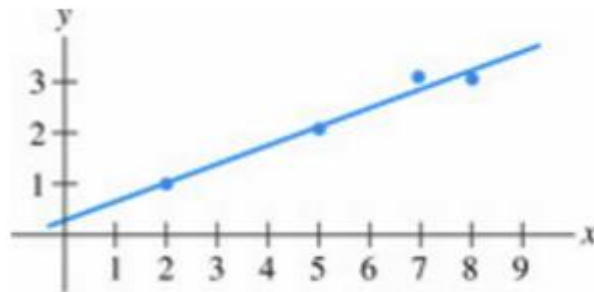


The vector \mathbf{b} is closer to $A\hat{\mathbf{x}}$ than to $A\mathbf{x}$ for other \mathbf{x} .

```
In [68]: err = norm(A*x - b)  
          print(err)  
  
          9.16515138991168
```


Example

- Find the equation $y = \beta_0 + \beta_1 x$ of the least-squares line that best fits the data points $(2, 1)$, $(5, 2)$, $(7, 3)$, and $(8, 3)$.



Solution in R and Python

```
> x = matrix(c(1,1,1,1,2,5,7,8), nrow = 4)
> y = matrix(c(1,2,3,3), nrow = 4)
> sol = ginv(x)%*%y
> print(sol)
      [,1]
[1,] 0.2857143
[2,] 0.3571429
```

```
In [70]: b = np.array([[1],[2],[3],[3]])
A = np.array([[1,2],[1,5],[1,7],[1,8]])
x = pinv(A)@b
print(x)

[[0.28571429]
 [0.35714286]]
```

How to read a csv file into R

```
myData = read.csv("U://MyDocuments/DataAnalytics/Stat2/Assignments/hw4_logistics/Oring.csv")
str(myData)
```

```
## 'data.frame':    24 obs. of  2 variables:
## $ Temp      : int  53 56 57 63 66 67 67 67 68 69 ...
## $ Failure: int   1 1 1 0 0 0 0 0 0 0 ...
```

```
summary(myData)
```

##	Temp	Failure
## Min.	:53.00	Min. :0.0000
## 1st Qu.:	67.00	1st Qu.:0.0000
## Median	:70.00	Median :0.0000
## Mean	:69.92	Mean :0.2917
## 3rd Qu.:	75.25	3rd Qu.:1.0000
## Max.	:81.00	Max. :1.0000

How to read a csv file into Python

```
In [81]: import pandas as pd
d=pd.read_csv('oring.csv')
d.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24 entries, 0 to 23
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Temp        24 non-null     int64
1    Failure     24 non-null     int64
dtypes: int64(2)
memory usage: 512.0 bytes
```

```
In [82]: d.describe()
```

Out[82]:

	Temp	Failure
count	24.000000	24.000000
mean	69.916667	0.291667
std	7.377502	0.464306
min	53.000000	0.000000
25%	67.000000	0.000000
50%	70.000000	0.000000
75%	75.250000	1.000000
max	81.000000	1.000000

Generating Random Numbers in R – rnorm and rep

```
myRand = rnorm(n = nrow(myData))
print(myRand)
```

```
## [1] -1.502733e+00 -6.416589e-01 -5.478076e-01 -4.825966e-01 -2.208273e-01
## [6] -1.057852e+00 -8.274779e-06 -6.115095e-01 -2.574924e-01 -1.702548e-01
## [11] -2.066829e+00 -1.034635e+00 -3.624287e-01 -2.228764e+00 1.800802e+00
## [16] 1.933479e-01 5.434829e-01 -1.779400e-01 -1.375305e+00 -9.575774e-01
## [21] -3.319702e-01 -3.919302e-01 -4.068536e-01 4.128651e-01
```

```
myZero = rep(0, nrow(myData))
print(myZero)
```

```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

Generating Random Numbers in Python – normal and repeat

```
In [99]: from numpy.random import normal
         np.random.seed(123)
         myRand = normal(size = d.shape[0]).reshape(d.shape[0],1)
         print(myRand)
```

```
[[-1.0856306 ]
 [ 0.99734545]
 [ 0.2829785 ]
 [-1.50629471]
 [-0.57860025]
 [ 1.65143654]
 [-2.42667924]
 [-0.42891263]
 [ 1.26593626]
 [-0.8667404 ]
 [-0.67888615]
 [-0.09470897]
 [ 1.49138963]
 [-0.638902 ]
 [-0.44398196]
 [-0.43435128]
 [ 2.20593008]
 [ 2.18678609]
 [ 1.0040539 ]
 [ 0.3861864 ]
 [ 0.73736858]
 [ 1.49073203]
 [-0.93583387]
 [ 1.17582904]]
```

```
: myZero = np.repeat(0, d.shape[0]).reshape(d.shape[0],1)
   print(myZero)
```

```
[[0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]
 [0]]
```

How to combine R arrays by columns (or rows)

```
myData = cbind(myZero, myRand, myData)
print(myData)
```

```
##      myZero      myRand Temp Failure
## 1      0 -1.502733e+00  53      1
## 2      0 -6.416589e-01  56      1
## 3      0 -5.478076e-01  57      1
## 4      0 -4.825966e-01  63      0
## 5      0 -2.208273e-01  66      0
## 6      0 -1.057852e+00  67      0
## 7      0 -8.274779e-06  67      0
## 8      0 -6.115095e-01  67      0
## 9      0 -2.574924e-01  68      0
## 10     0 -1.702548e-01  69      0
## 11     0 -2.066829e+00  70      0
## 12     0 -1.034635e+00  70      1
## 13     0 -3.624287e-01  70      1
## 14     0 -2.228764e+00  70      1
## 15     0  1.800802e+00  72      0
## 16     0  1.933479e-01  73      0
## 17     0  5.434829e-01  75      0
## 18     0 -1.779400e-01  75      1
## 19     0 -1.375305e+00  76      0
## 20     0 -9.575774e-01  76      0
## 21     0 -3.319702e-01  78      0
## 22     0 -3.919302e-01  79      0
## 23     0 -4.068536e-01  80      0
## 24     0  4.128651e-01  81      0
```

How to combine Python arrays by columns (or rows)

```
In [103]: np.hstack([myZero, myRand, d])
```

```
Out[103]: array([[ 0.          , -1.0856306 , 53.          , 1.          ],
 [ 0.          ,  0.99734545, 56.          , 1.          ],
 [ 0.          ,  0.2829785 , 57.          , 1.          ],
 [ 0.          , -1.50629471, 63.          , 0.          ],
 [ 0.          , -0.57860025, 66.          , 0.          ],
 [ 0.          ,  1.65143654, 67.          , 0.          ],
 [ 0.          , -2.42667924, 67.          , 0.          ],
 [ 0.          , -0.42891263, 67.          , 0.          ],
 [ 0.          ,  1.26593626, 68.          , 0.          ],
 [ 0.          , -0.8667404 , 69.          , 0.          ],
 [ 0.          , -0.67888615, 70.          , 0.          ],
 [ 0.          , -0.09470897, 70.          , 1.          ],
 [ 0.          ,  1.49138963, 70.          , 1.          ],
 [ 0.          , -0.638902  , 70.          , 1.          ],
 [ 0.          , -0.44398196, 72.          , 0.          ],
 [ 0.          , -0.43435128, 73.          , 0.          ],
 [ 0.          ,  2.20593008, 75.          , 0.          ],
 [ 0.          ,  2.18678609, 75.          , 1.          ],
 [ 0.          ,  1.0040539 , 76.          , 0.          ],
 [ 0.          ,  0.3861864 , 76.          , 0.          ],
 [ 0.          ,  0.73736858, 78.          , 0.          ],
 [ 0.          ,  1.49073203, 79.          , 0.          ],
 [ 0.          , -0.93583387, 80.          , 0.          ],
 [ 0.          ,  1.17582904, 81.          , 0.          ]])
```


Reference

- Linear Algebra and Its Applications by David Lay